


# **NEURAL NETWORKS I**

**DATA/MSML 603: Principles of Machine Learning**

# Neural Function



- Brain function (thought) occurs as the result of the firing of **neurons**
- Neurons connect to each other through **synapses**, which propagate **action potential** (electrical impulses) by releasing **neurotransmitters**
  - Synapses can be excitatory (potential-increasing) or inhibitory (potential-decreasing), and have varying activation thresholds
  - Learning occurs as a result of the synapses' plasticity: They exhibit long-term changes in connection strength (causes structural and functional change in nervous system)
- There are about  $10^{11}$  neurons and about  $10^{14}$  synapses in the human brain!

# Biology of A Neuron






Image source: Bruce Blausen



# Brain Structure



- Different areas of the brain have different functions
  - Some areas seem to have the same function in all humans (e.g., Broca's region for motor speech); the overall layout is generally consistent
  - Some areas are more plastic, and vary in their function; also, the lower-level structure and function vary greatly
- We do not know how different functions are “assigned” or acquired
  - Partly the result of the physical layout / connection to inputs (sensors) and outputs (effectors)
  - Partly the result of experience (learning)

# Comparison of Computing Power

INFORMATION CIRCA 2012	Computer	Human Brain
Computation Units	10-core Xeon: $10^9$ Gates	$10^{11}$ Neurons
Storage Units	$10^9$ bits RAM, $10^{12}$ bits disk	$10^{11}$ neurons, $10^{14}$ synapses
Cycle time	$10^{-9}$ sec	$10^{-3}$ sec
Bandwidth	$10^9$ bits/sec	$10^{14}$ bits/sec


- (Artificial) neural networks are made up of **nodes** or **units**, connected by **links**
- Each link has an associated **weight** and **activation level**
- Each node has an **input function** (typically summing over weighted inputs), an **activation function**, and an **output**

# (Artificial) Neural Networks


- Origins: Algorithms that try to mimic the brain
- Very widely used in 80s and early 90s; popularity diminished in late 90s
- Recent resurgence: State-of-the-art technique for many applications
- Artificial neural networks are not nearly as complex or intricate as the actual brain structure

# Neural Networks (2/2)


- Computers are way faster than neurons...
- But, there are a lot more neurons than we can reasonably model in modern digital computers, and they all fire in parallel
- Neural networks are designed to be massively parallel
- The brain is effectively a billion times faster



# Neuron Model: Threshold Unit




# Neuron Model: Logistic Unit



$$\begin{aligned} h_{\boldsymbol{\theta}}(\mathbf{x}) &= g(\boldsymbol{\theta}^T \mathbf{x}) \\ &= \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}}} \end{aligned}$$


Sigmoid (logistic) activation function:  $g(z) = \frac{1}{1 + e^{-z}}$



# Neuron Model: Other Activation Functions

- Tanh or hyperbolic tangent activation function
- Rectified Linear Unit (ReLU) activation function
- Leaky ReLU
- We will come back to these later...

# Neural Network




# Feed-forward Process (1/2)


- Input layer units are set by some exterior function (think of these as **sensors**), which causes their output links to be **activated** at the specified level
- Working forward through the network, the **input function** of each unit is applied to compute the input value
  - Usually this is just the weighted sum of the activation on the links feeding into this node

# Feed-forward Process (2/2)

- The **activation function** transforms this input function (value) into a final value
  - Typically this is a **nonlinear** function, often a sigmoid function corresponding to the threshold of that node



# Neural Network



$$\begin{aligned} a_1^{(2)} &= g(\Theta_{10}^{(1)}x_0 + \Theta_{11}^{(1)}x_1 + \Theta_{12}^{(1)}x_2 + \Theta_{13}^{(1)}x_3) \\ a_2^{(2)} &= g(\Theta_{20}^{(1)}x_0 + \Theta_{21}^{(1)}x_1 + \Theta_{22}^{(1)}x_2 + \Theta_{23}^{(1)}x_3) \\ a_3^{(2)} &= g(\Theta_{30}^{(1)}x_0 + \Theta_{31}^{(1)}x_1 + \Theta_{32}^{(1)}x_2 + \Theta_{33}^{(1)}x_3) \\ h_{\theta}(x) &= a_1^{(3)} = g(\Theta_{10}^{(2)}a_0^{(2)} + \Theta_{11}^{(2)}a_1^{(2)} + \Theta_{12}^{(2)}a_2^{(2)} + \Theta_{13}^{(2)}a_3^{(2)}) \end{aligned}$$

If network has  $s_j$  units in layer  $j$  and  $s_{j+1}$  units in layer  $j+1$ ,  
then  $\Theta^{(j)}$  has dimension  $s_{j+1} \times (s_j + 1)$  [includes bias].

}  
Hidden units

$$\Theta^{(1)} \in \mathbb{R}^{3 \times 4} \quad \Theta^{(2)} \in \mathbb{R}^{1 \times 4}$$



# Vectorization

$$a_1^{(2)} = g\left(\Theta_{10}^{(1)}x_0 + \Theta_{11}^{(1)}x_1 + \Theta_{12}^{(1)}x_2 + \Theta_{13}^{(1)}x_3\right) = g\left(z_1^{(2)}\right)$$

$$a_2^{(2)} = g\left(\Theta_{20}^{(1)}x_0 + \Theta_{21}^{(1)}x_1 + \Theta_{22}^{(1)}x_2 + \Theta_{23}^{(1)}x_3\right) = g\left(z_2^{(2)}\right)$$

$$a_3^{(2)} = g\left(\Theta_{30}^{(1)}x_0 + \Theta_{31}^{(1)}x_1 + \Theta_{32}^{(1)}x_2 + \Theta_{33}^{(1)}x_3\right) = g\left(z_3^{(2)}\right)$$

$$h_{\Theta}(\mathbf{x}) = g\left(\Theta_{10}^{(2)}a_0^{(2)} + \Theta_{11}^{(2)}a_1^{(2)} + \Theta_{12}^{(2)}a_2^{(2)} + \Theta_{13}^{(2)}a_3^{(2)}\right) = g\left(z_1^{(3)}\right)$$



Feed-Forward Steps:

$$\mathbf{z}^{(2)} = \Theta^{(1)}\mathbf{x}$$

$$\mathbf{a}^{(2)} = g(\mathbf{z}^{(2)})$$

$$\text{Add } a_0^{(2)} = 1$$

$$\mathbf{z}^{(3)} = \Theta^{(2)}\mathbf{a}^{(2)}$$


$$h_{\Theta}(\mathbf{x}) = \mathbf{a}^{(3)} = g(\mathbf{z}^{(3)})$$

bias



SCIENCE  
ACADEMY

# Other Network Architectures



- $L$  denotes the number of layers
- $s \in \mathbb{N}^{+L}$  contains the numbers of nodes at each layer
  - Not counting bias units
  - Typically,  $s_0 = d$  (# input features) and  $s_{L-1} = K$  (# classes)

# Multiple Output Units: One-vs-Rest



Pedestrian




Car



Motorcycle



Truck



$$h_{\Theta}(\mathbf{x}) \in \mathbb{R}^K$$

We want:


$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix} \quad \text{when pedestrian}$$

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix} \quad \text{when car}$$

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \quad \text{when motorcycle}$$

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} \quad \text{when truck}$$


# Multiple Output Units: One-vs-Rest



- Given  $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$
- Must convert labels to 1-of- $K$  representation

– e.g.,  $y_i = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$  when motorcycle,  $y_i = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$  when car, etc.

# Neural Network Classification



Binary classification

$y = 0 \text{ or } 1$

1 output unit ( $s_{L-1} = 1$ )

**Given:**

$$\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$$

$\mathbf{s} \in \mathbb{N}^{+L}$  contains # nodes at each layer

–  $s_0 = d$  (# features)

Multi-class classification ( $K$  classes)

$$\mathbf{y} \in \mathbb{R}^K \quad \text{e.g. } \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

pedestrian   car   motorcycle   truck

$K$  output units ( $s_{L-1} = K$ )


# Understanding Representations

# Representing Boolean Functions


Simple example: AND

$$x_1, x_2 \in \{0, 1\}$$

$$y = x_1 \text{ AND } x_2$$



$$\begin{aligned} h_{\theta}(\mathbf{x}) &= g(-30 + 20x_1 + 20x_2) \\ &= \frac{1}{1 + e^{-(30+20x_1+20x_2)}} \end{aligned}$$







$x_1$	$x_2$	$h_{\theta}(\mathbf{x})$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$


$$g(-30) = 9.36 \cdot 10^{-14}, \quad g(-10) = 4.54 \cdot 10^{-5}$$

$$g(-30) = 9.36 \cdot 10^{-14}, \quad g(-10) = 4.54 \cdot 10^{-5}$$


# Representing Boolean Functions



# Combining Representations





# Layering Representations



$x_1 \dots x_{20}$   
 $x_{21} \dots x_{40}$   
 $x_{41} \dots x_{60}$   
•  
•  
•  
 $x_{381} \dots x_{400}$

# Layering Representations (2/2)



Visualization of Hidden Layer


# Convolutional Neural Networks

# Convolutional Neural Networks (CNNs)

- Originally introduced for problems where input data has a grid-like structure
  - Image – pixels reside on a two-dimensional grid
  - Audio waveform – one-dimensional grid
  - Volumetric data, e.g., computer tomography (CT) scans or video data – three-dimensional grid


# Convolutional Neural Networks (CNNs)

- Putting all input variables representing the image pixels in a long vector causes loss of a lot of structure present in the image
  - For instance, two pixels close to each other typically have more in common than two pixels farther apart, which is destroyed by vectorization



# Convolutional Layer

- CNNs tries to preserve this information by representing both input variables and hidden layers as matrices



**Dense layer:** each input variable connected to all hidden units in subsequent layer, and each such connection has a unique parameter associated to it

- Provide too much flexibility for images

**Convolutional layer** exploits the structure present in images to find a more efficiently parameterized model

# Convolutional Layer

- Convolutional layer leverages two key concepts
  1. Sparse interactions: Most of the parameters in a corresponding dense layer are forced to be zero
    - A hidden unit in a convolutional layer only depends on the pixels in a small region of the image and not on all pixels
  2. Parameter sharing: In a dense layer, each link between an input variable and a hidden unit has its own unique parameter
    - In a convolutional layer, the set of parameters (filters), for different hidden units are all the same
    - Instead of learning separate sets of parameters, learn only one set of parameters

# Convolutional Layer

- Sparse interactions and parameter sharing in a convolutional layer makes CNNs relatively invariant to translations of objects in images
  - If the parameters in a filter are sensitive to a certain detail (e.g., a corner, an edge), a hidden unit will react to this detail regardless of where in the image that detail is present
- A convolutional layer uses significantly fewer parameters compared to a dense layer

# Convolution


- Loosely mimics neuronal responses
- To compute:
  - Slide small kernel (also called a filter) over input
  - Element-wise multiplication and sum

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input matrix

1	0	1
0	1	0
1	0	1

2D Convolutional  
3x3 filter



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		


Convolved  
Feature

3	0	3	1	2	2	1	0
0	2	0	2	1	0	3	1
3	0	1	1	2	2	2	3
2	0	0	0	2	2	2	2
2	0	0	0	0	0	1	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0
2	0	0


# Convolution Summary

- Layers structured as matrices, not vectors
- Convolution of matrices by small (e.g. 3 x 3) matrices called **kernels**
- The kernels are the **learned weights**
- Depending on **padding** and **stride**, image may decrease slightly in size
- Intuitively, convolutional layers learn how to relate neighboring pixels
- Image layers in a CNN typically have many **features/hidden units**
  - Organized as “**channels**”



# Motivation: Conv. for Edge Detection

$\begin{bmatrix} -1 & -1 & -1 \\ 2 & 2 & 2 \\ -1 & -1 & -1 \end{bmatrix}$	$\begin{bmatrix} -1 & 2 & -1 \\ -1 & 2 & -1 \\ -1 & 2 & -1 \end{bmatrix}$
Horizontal lines	Vertical lines
$\begin{bmatrix} -1 & -1 & 2 \\ -1 & 2 & -1 \\ 2 & -1 & -1 \end{bmatrix}$	$\begin{bmatrix} 2 & -1 & -1 \\ -1 & 2 & -1 \\ -1 & -1 & 2 \end{bmatrix}$
45 degree lines	135 degree lines




Horizontal lines

# Cross-correlation in Two Dimensions

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0


Pixel representation of filter



Visualization of a curve detector filter



Original image



Visualization of the filter on the image



Visualization of the receptive field

0	0	0	0	0	0	30	0
0	0	0	0	50	50	50	
0	0	0	20	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	
0	0	0	50	50	0	0	

Pixel representation of the receptive field

\*


0	0	0	0	0	0	30	0
0	0	0	0	30	0	0	
0	0	0	30	0	0	0	
0	0	0	30	0	0	0	
0	0	0	30	0	0	0	
0	0	0	30	0	0	0	
0	0	0	0	0	0	0	
0	0	0	0	0	0	0	

Pixel representation of filter

$$\text{Multiplication and Summation} = (50*30)+(50*30)+(50*30)+(20*30)+(50*30) = 6600 \text{ (A large number!)}$$

# Convolutional Layer with Strides

- With zero-padding, in a convolutional layer we saw earlier, the number of hidden units is equal to the number of pixels inn an image
- We often want to reduce the number of hidden units and store only the most important information from the previous layers
- Stride controls how many pixels the filter/kernel shifts over at each step



[Source: <https://svitla.com/blog/math-at-the-heart-of-cnn/>]


© 2020, UNIVERSITY OF MARYLAND

# Pooling Layer

- Another way of summarizing information in previous layers
- Additional layer after a convolutional layer
  - Similar to convolutional layer, depends only on a small region of pixels
  - Does not require extra trainable parameters
  - Also use strides to condense the information
- Most common pooling methods
  - Max pooling (MaxPool)
  - Average pooling


# Max Pool (1/2)

- Intuitively, MaxPool is used because nearby pixels will give very similar output
- Over the course of a network, pooling shifts layers from spatial dimensions to feature



# Max Pool (2/2)

- Practice




# Channels


- Using a single filter/kernel in each convolutional layer often not sufficient to encode all interesting properties of images
- Add multiple filters, each with own set of parameters to be trained
  - Each filter produces its own set of hidden units, called a “channel”, using the same convolution operation
  - Each layer of hidden units in a CNN organized into a tensor with dimension (rows x columns x channels)
- When stacking convolutional layers, each filter depends on all the channels in the previous layer
  - Each filter is a tensor of dimension (filter rows x filter columns x input channels)

# Color Images

- Color (RGB) image represented as a tensor ( $h \times w \times 3$ )
- 3D filters
  - Depth of each filter chosen to match the number of color channels
  - Each channel convolved with own 2D filter
  - Add the values from the convolution with 2D filters
- Each 3D filter produces a separate 2D matrix with hidden units




# What CNNs Learn




Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

# Deep Belief Net on Face Images



# Learning of Object Parts




SCIENCE  
ACADEMY

© 2020, UNIVERSITY OF MARYLAND

DATA/MSML 603

# Training on Multiple Objects



Trained on 4 classes (cars, faces, motorbikes, airplanes).

Second layer: Shared-features and object-specific features.

Third layer: More specific features.