

# **NEURAL NETWORKS – PART 3**

**DATA/MSML 603: Principles of Machine Learning**

# Last Time

- Training Neural Networks
  - Stochastic gradient descent
  - Backpropagation – forward propagation & backward propagation
- Activation functions
  - Vanishing gradient & exploding gradient
- Loss functions
- Underfitting & Overfitting

# Ensemble Methods

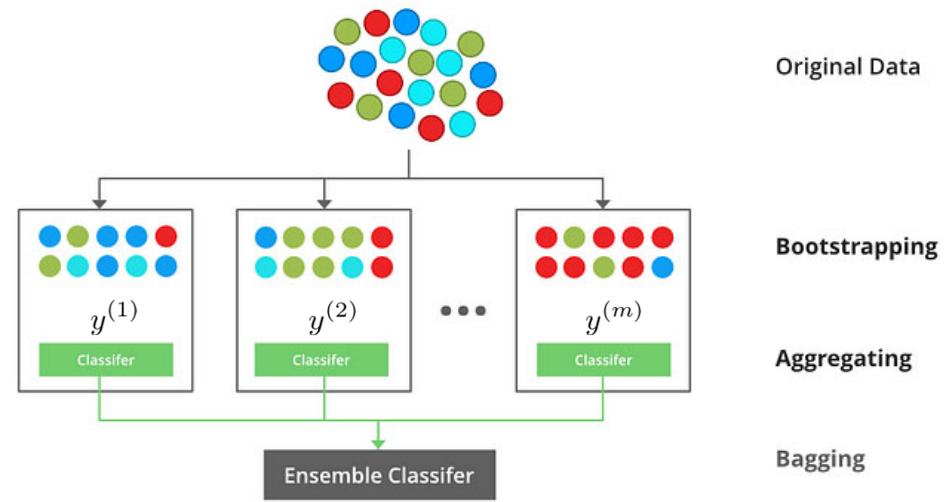
# Bootstrap Aggregating (Bagging)

- Recall bias-variance trade-off
  - More flexible a model, the lower its bias will be and the higher the risk of overfitting (i.e., higher model variance)
- Model variance of flexible models can be reduced by using them as base (or ensemble) models in bootstrap aggregating (bagging) without increasing bias
- Bootstrap – method for artificially creating multiple datasets from one dataset
  - All datasets have the same size (as the original dataset)
  - Intuition: Suppose the original dataset is representative of real samples/distribution, i.e., if we collected more training data, the new data points/samples would likely be similar to the existing training data. Then, randomly choosing samples from the original dataset would be a reasonable way to simulate new training datasets
  - Each bootstrapped dataset obtained using **sampling with replacement**
    - May contain more than one copy of some of the original training data points/samples

# Bootstrap Aggregating (Bagging)

- Bootstrap Algorithm used to generate  $m$  datasets
- Multiple ( $m$ ) classifiers are trained in parallel using  $m$  datasets
- Output of ensemble classifier given by average

$$\hat{y}_{bag}(\mathbf{x}) = \frac{1}{m} \sum_{l=1}^m \hat{y}^{(l)}(\mathbf{x})$$



[source: medium.com]

$$\hat{y}_{bag}(\mathbf{x}) = \frac{1}{m} \sum_{l=1}^m \hat{y}^{(l)}(\mathbf{x})$$

# Bootstrap Aggregating (Bagging)

- But, how does it reduce model variance?
- Suppose  $\hat{y}^{(1)}(\mathbf{x}), \hat{y}^{(2)}(\mathbf{x}), \dots, \hat{y}^{(m)}(\mathbf{x})$  denote the predictions from individual ensemble members
- We can view them as identically distributed random variables, but not necessarily independent with mean  $\mathbb{E}[\hat{y}^{(1)}(\mathbf{x})] = \mu$  and  $Var(\hat{y}^{(1)}(\mathbf{x})) = \sigma^2$

$$\mathbb{E} [\hat{y}_{bag}(\mathbf{x})] = \mathbb{E} \left[ \frac{1}{m} \sum_{l=1}^m \hat{y}^{(l)}(\mathbf{x}) \right] = \mu$$

$$Var \left( \frac{1}{m} \sum_{l=1}^m \hat{y}^{(l)}(\mathbf{x}) \right) = \frac{1 - \rho}{m} \sigma^2 + \rho \sigma^2 < \sigma^2 \quad \text{if } \rho < 1$$

$\rho$ — Correlation coefficient between predictions of two ensemble members

# Boosting

- Primarily used for **reducing bias** in high-bias base models (often simple models)
  - Reduce the bias by turning ensemble of simple/weak models into one strong model
- Unlike in bagging where multiple models are trained in parallel, boosting uses a **sequential construction of ensemble members** (cannot be done in parallel)
  - Each model tries to correct the mistakes of previous model by modifying training dataset at each iteration – put larger weights on data points/samples for which the previous models performed poorly
- Two key details that need to be determined
  - How to compute the weights for training data points/samples at each iteration
  - How to combine ensemble models to construct the final model

# AdaBoost (Adaptive Boosting)

- First successful implementation of boosting
- Hard prediction for binary classification:  $\hat{y}^{(1)}(\mathbf{x}), \hat{y}^{(2)}(\mathbf{x}), \dots, \hat{y}^{(m)}(\mathbf{x})$
- Boosted classifier using a **weighted** majority vote

$$\hat{y}_{\text{boost}}^{(m)}(\mathbf{x}) = \text{sign} \left\{ \sum_{l=1}^m \alpha^{(l)} \hat{y}^{(l)}(\mathbf{x}) \right\}$$

- Each ensemble member votes either -1 or +1
- Coefficients  $\alpha^{(l)} > 0$  can be viewed as a level of **confidence** in the prediction by  $l$ -th ensemble member

# AdaBoost (Adaptive Boosting)

- Ensemble members and their coefficients  $\alpha^{(l)}$  are trained greedily by minimizing **exponential loss** of boosted classifier at each iteration
- For each  $l = 0, 1, \dots, m$ , define  $f^{(l)}(\mathbf{x}) = \sum_{k=1}^l \alpha^{(k)} \hat{y}^{(k)}(\mathbf{x})$  with  $f^{(0)}(\mathbf{x}) = 0$
- Boosted classifier after  $l$  iterations  $\hat{y}_{boost}^{(l)}(\mathbf{x}) = \text{sign}\{f^{(l)}(\mathbf{x})\}$
- $\hat{y}^{(l)}$  is selected to minimize the exponential loss

$$\begin{aligned} \sum_{i=1}^n L(y^i \cdot f^{(l)}(\mathbf{x}^i)) &= \sum_{i=1}^n \exp \left( -y^i \left( \underbrace{f^{(l-1)}(\mathbf{x}^i)}_{+} + \boxed{\alpha^{(l)} \hat{y}^{(l)}(\mathbf{x}^i)} \right) \right) \\ &= f^{(l)}(\mathbf{x}^i) \end{aligned}$$

# AdaBoost (Adaptive Boosting)

- At the  $l$ -th iteration, need to find

$$\begin{aligned}(\alpha^{(l)}, \hat{y}^{(l)}) &= \arg \min_{(\alpha, \hat{y})} \sum_{i=1}^n \exp \left( -y^i (f^{(l-1)}(\mathbf{x}^i) + \alpha \hat{y}(\mathbf{x}^i)) \right) & (\alpha > 0) \\&= \arg \min_{(\alpha, \hat{y})} \sum_{i=1}^n \underbrace{\exp \left( -y^i f^{(l-1)}(\mathbf{x}^i) \right)}_{=: w_i^{(l)}} \exp \left( -y^i \alpha \hat{y}(\mathbf{x}^i) \right) & (e^{a+b} = e^a \cdot e^b) \\&= \arg \min_{(\alpha, \hat{y})} \sum_{i=1}^n w_i^{(l)} \exp \left( -y^i \alpha \hat{y}(\mathbf{x}^i) \right)\end{aligned}$$

- Define  $W = \sum_{i=1}^n w_i^{(l)}$  be the total weight sum

# AdaBoost (Adaptive Boosting)

- Break the function to be minimized into two terms – one for data points that are correctly classified by candidate  $\hat{y}$  and the other for incorrectly classified data points

$$\sum_{i=1}^n w_i^{(l)} \exp(-y^i \alpha \hat{y}(\mathbf{x}^i)) = e^{-\alpha} \underbrace{\sum_{i=1}^n w_i^{(l)} \mathbf{1}\{y^i = \hat{y}(\mathbf{x}^i)\}}_{=: W_c} + e^{\alpha} \underbrace{\sum_{i=1}^n w_i^{(l)} \mathbf{1}\{y^i \neq \hat{y}(\mathbf{x}^i)\}}_{=: W_e}$$
$$(W = W_c + W_e) = e^{-\alpha} W + (e^{\alpha} - e^{-\alpha}) W_e > 0$$

Weighted  
misclassification  
loss

- If we want to minimize the exponential loss, we need to minimize  $W_e$ , i.e.,

$$\hat{y}^{(l)} = \arg \min_{\hat{y}} \sum_{i=1}^n w_i^{(l)} \mathbf{1}\{y^i \neq \hat{y}(\mathbf{x}^i)\}$$

# AdaBoost (Adaptive Boosting)

$$\hat{y}^{(l)} = \arg \min_{\hat{y}} \sum_{i=1}^n w_i^{(l)} \mathbf{1} \{y^i \neq \hat{y}(\mathbf{x}^i)\}$$

- The  $l$ -th ensemble member should be trained by minimizing the **weighted misclassification loss** with weights  $w_i^{(l)} = \exp(-y^i f^{(l-1)}(\mathbf{x}^i))$ 
  - At iteration  $l$ , should focus our attention on the data points/samples previously misclassified to correct the mistakes of the first  $l-1$  classifiers
- What about the coefficient  $\alpha^{(l)}$ ?

# AdaBoost (Adaptive Boosting)

- Recall that the function we want to minimize is given by

$$\sum_{i=1}^n w_i^{(l)} \exp(-y^i \alpha \hat{y}(\mathbf{x}^i)) = e^{-\alpha} + (e^\alpha - e^{-\alpha})W_e \quad (10-1)$$

- Once we choose  $\hat{y}^{(l)}$ ,  $W_e = \sum_{i=1}^n w_i^{(l)} \mathbf{1}\{y^i \neq \hat{y}^{(l)}(\mathbf{x}^i)\}$  is a fixed constant

- Differentiate (10-1) w.r.t.  $\alpha$  and set the derivative to zero

$$-e^{-\alpha}W + W_e(e^\alpha - (-1)e^{-\alpha}) = -e^{-\alpha}W + W_e(e^\alpha + e^{-\alpha}) = 0$$

$$\iff W = W_e(e^{2\alpha} + 1)$$

$$\iff \alpha^{(l)} = \frac{1}{2} \log\left(\frac{W}{W_e} - 1\right)$$

Reflects the confidence  
in the classifier's  
predictions

Depends on the training error of  $\hat{y}^{(l)}$

# Recurrent Neural Networks

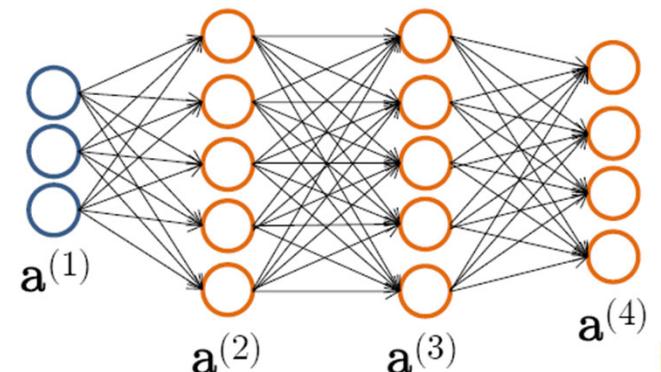
# Feedforward NN

- Neural network used to approximate the relationship between the input (e.g., feature values) and the output (e.g., label)

- For each example, input values determine output value, **independently of other examples**

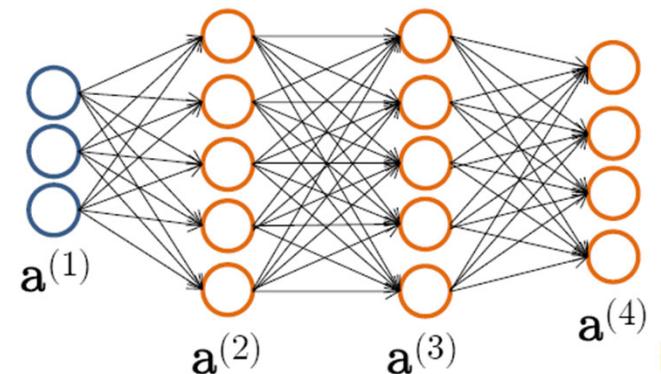
$$\mathbf{y} = f(\mathbf{x})$$

- Universal function approximator – when the neural network size is sufficiently large, it can approximate a large class of functions using appropriate weights



# Feedforward NN

- In some cases, inputs are sequential and outputs cannot be computed independently
  - Often need to discover the pattern in sequential data (e.g., time series forecasting)
  - Requires some “memory” of the past
- In addition, a neural network is trained using fixed-size inputs, i.e., the dimension of inputs is pre-determined and cannot be changed on the fly
- Question: What if we don’t know the exact length of the inputs that need to be handled by a neural network or the inputs are sequential data of variable length?



# Example

- Predicting the next word in the sequence
  - *He was born in Germany and grew up in Hamburg till he was a teenager. It is no wonder that he is fluent in .....*
  - *Tom and Mary are married, and Tom is her .....*
- Machine translation
  - *Come arrivare al Duomo di Milano?      ➡      How do I get to Duomo di Milano?*
- Time series prediction
- Sentiment analysis

# Recurrent Neural Networks

- Dates back to (Rumelhart *et al.*, 1986)
- A family of neural networks for handling **sequential data**, which involves variable length inputs or outputs

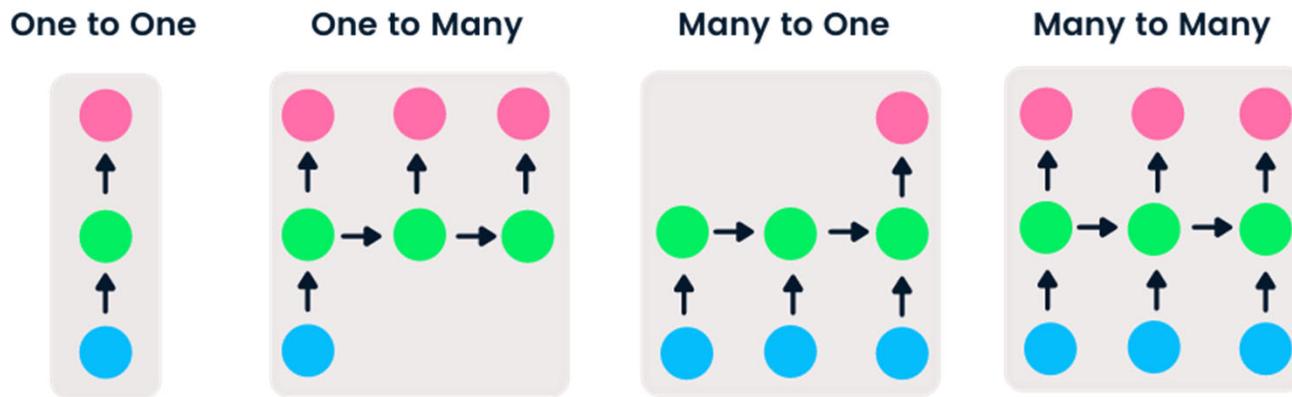
# Sequential Data

- Each data point: A sequence of vectors  $x(t)$ , for  $1 \leq t \leq \tau$
- Batch data: many sequences with different lengths  $\tau$
- Label: can be a scalar, a vector, or even a sequence

# Types of RNNs

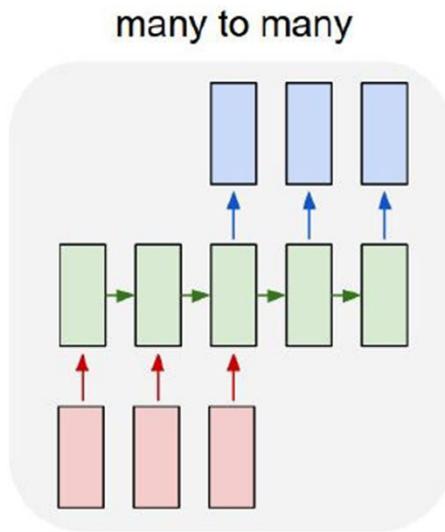
- Depending on the lengths/number of inputs and outputs
  - One-to-one – single input and single output
  - One-to-many – single input and multiple outputs
    - Generating image captions
  - Many-to-one – takes a sequence of multiple inputs and predicts a single output
    - Sentiment classification (e.g., input = text with multiple words, output = category)
  - Many-to-many – takes multiple inputs and produces multiple outputs
    - Machine language translation

# Types of RNNs

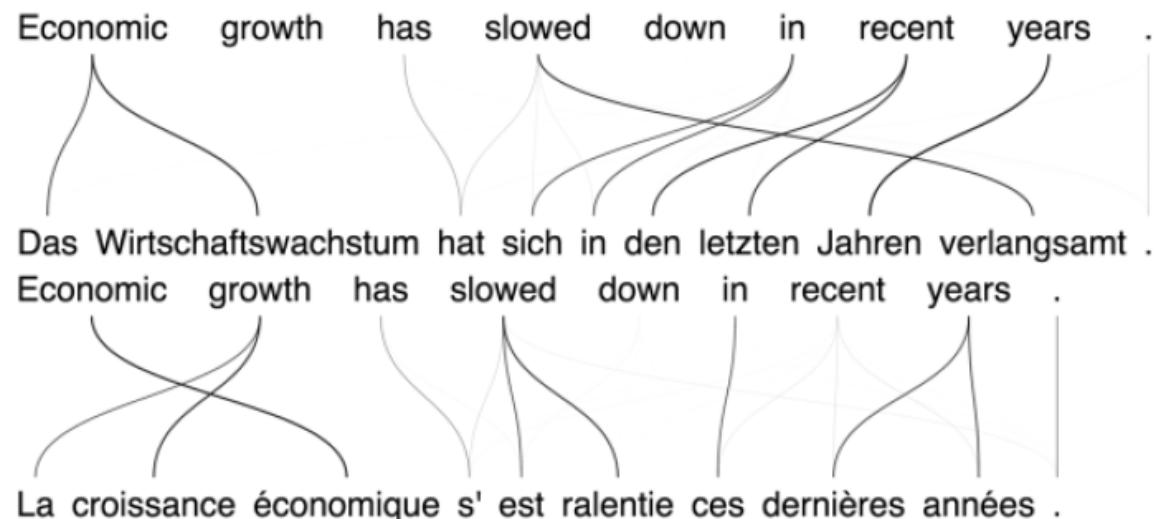


[source: [datacamp.com](https://www.datacamp.com)]

# Sequential Data: Machine Translation



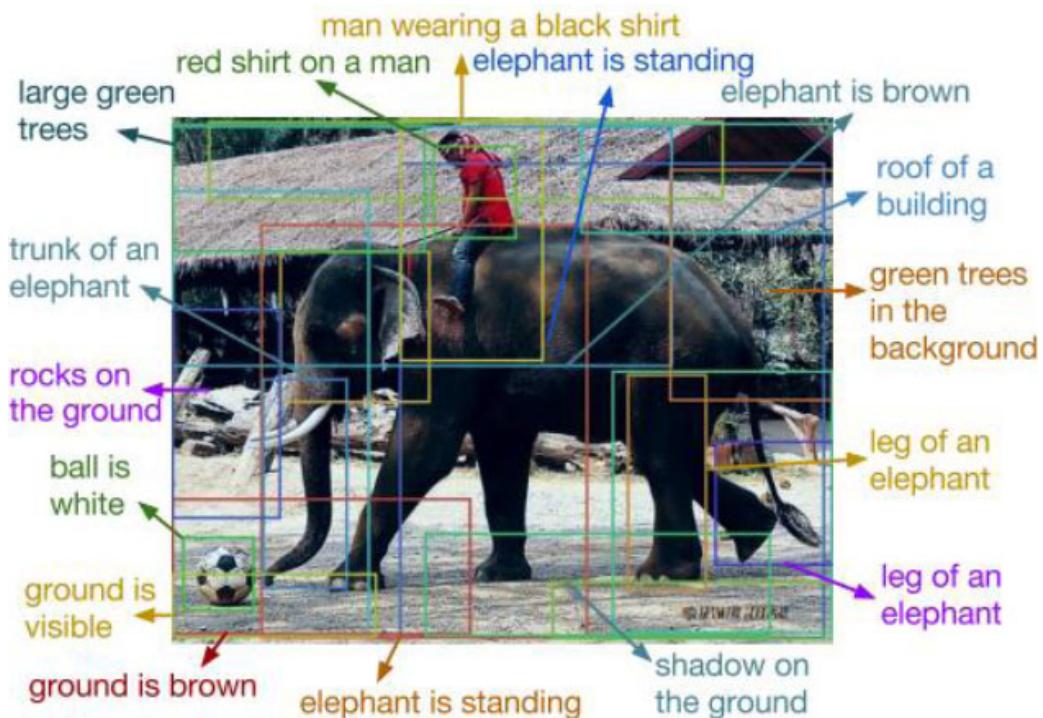
e.g. **Machine Translation**  
seq of words -> seq of words



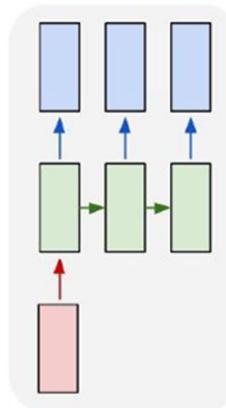
# More Complicated Sequential Data

- **Data point:** two dimensional sequences like images
- **Label:** different type of sequences like text sentences
- Example: image captioning

# Image Captioning



one to many



e.g. **Image Captioning**  
image -> sequence of words

Image source: "DenseCap: Fully Convolutional Localization Networks for Dense Captioning," by Justin Johnson, Andrej Karpathy, Li Fei-Fei

© 2020, UNIVERSITY OF MARYLAND

# Key Features

- Can handle inputs of variable lengths with a **fixed number of parameters**
- Parameter sharing: **same weights used for all inputs** (reduced complexity for generalization)
- **Hidden state** or **memory state** (computed for each input)
  - Used to keep memory of past inputs
  - Evolves over time steps
  - Weights updated using backpropagation through time (BPTT)

# A System Driven By External Data

- Discrete-time dynamical system driven by external input
  - $s^{(t)}$  : state of the system at time t
  - $x^{(t)}$  : input at time t
- State of the system at time  $t+1$  depends on two things
  - State of the system at time t
  - Input at time t

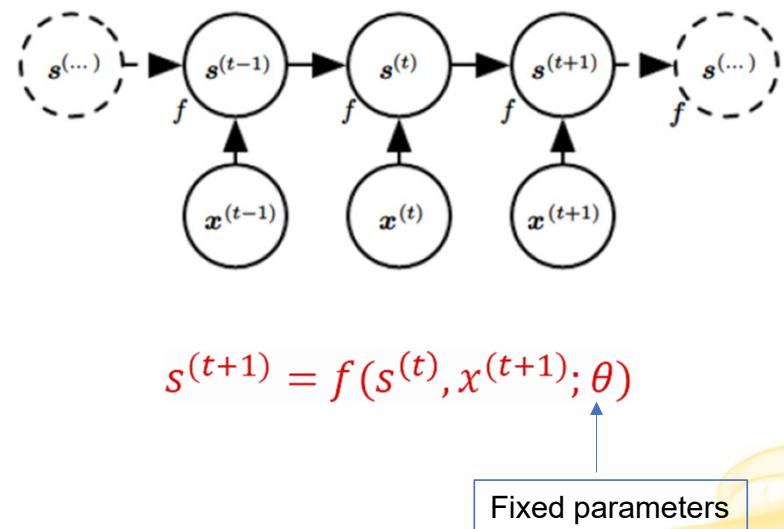
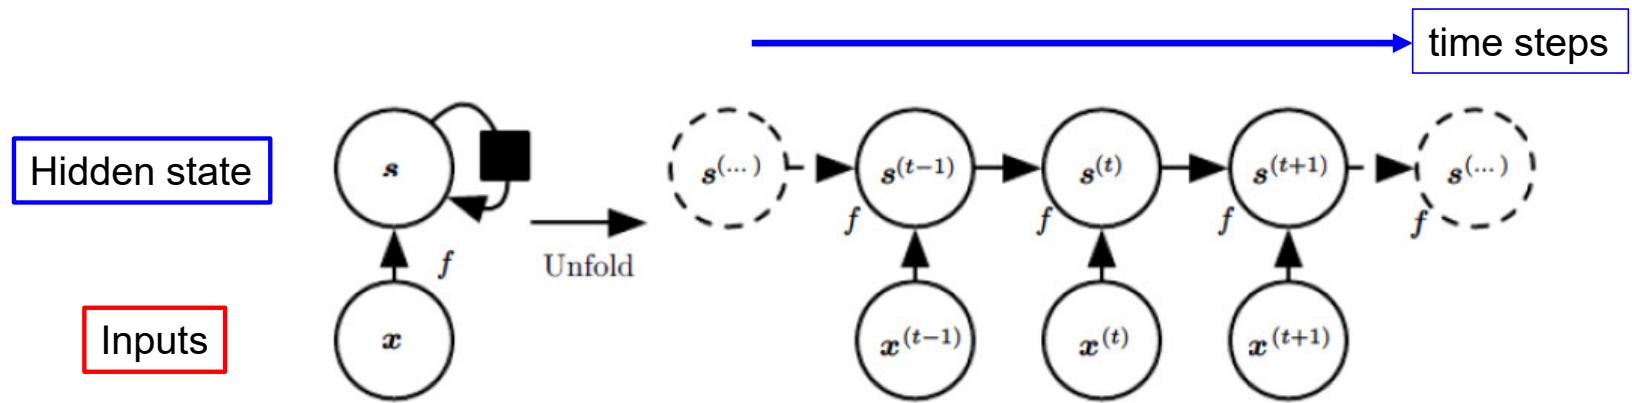


Figure from Deep Learning, Goodfellow, Bengio and Courville

# Compact View



$$s^{(t+1)} = f(s^{(t)}, x^{(t+1)}; \theta)$$

Figure from Deep Learning, Goodfellow, Bengio and Courville

# Compact View

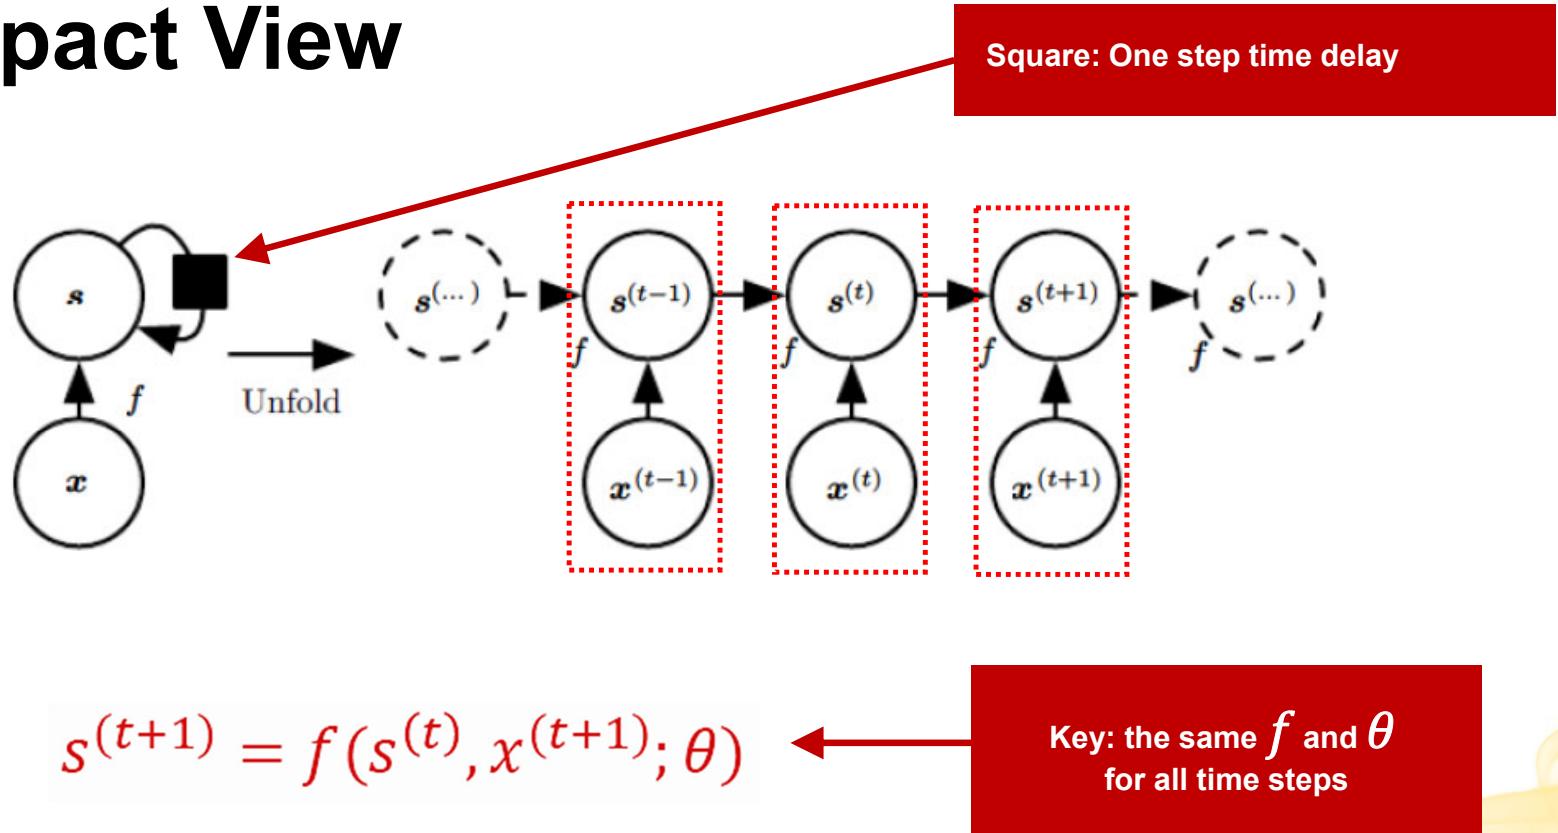


Figure from Deep Learning, Goodfellow, Bengio and Courville

# Recurrent Neural Networks

- Use **the same computational function ( $f$ ) and parameters ( $\theta$ )** across different time steps of the sequence (for better generalization)
- Each time step: takes the input entry and **the previous hidden state** to compute the output entry
- Loss: typically computed at every time step

# Recurrent Neural Networks (single layer)

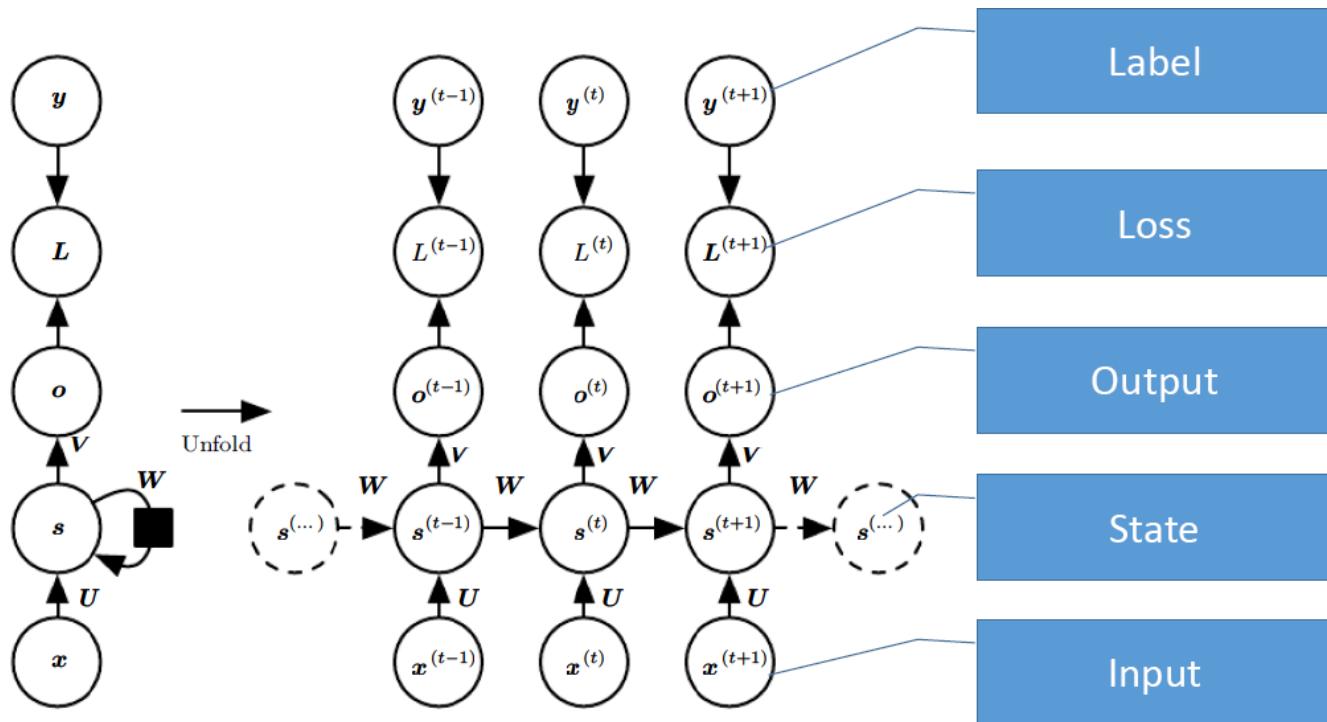


Figure from Deep Learning, Goodfellow, Bengio and Courville

# Recurrent Neural Networks

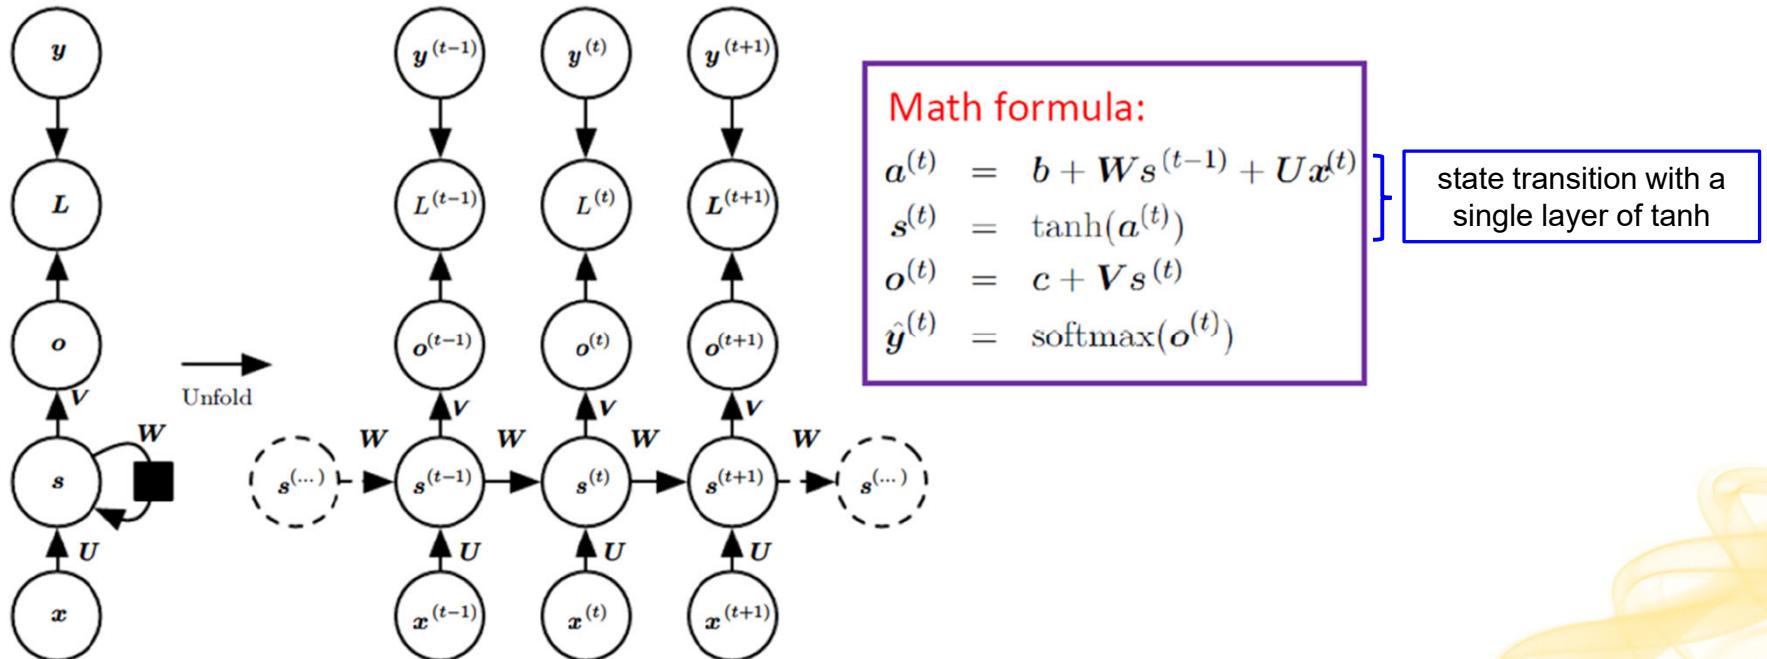


Figure from Deep Learning, Goodfellow, Bengio and Courville

# Advantages

- Hidden state: lossy summary of the past
- Shared functions and parameters: greatly reduce the capacity and good for **generalization** in learning
- Explicitly use the **prior knowledge** that the sequential data can be processed in the same way at different time step (e.g., NLP)
- Yet still powerful (Turing-complete): any function computable by a Turing machine can be computed by such a recurrent network of a finite size (see, e.g., Siegelmann and Sontag [1995])

# Training RNN

- Principle: unfold the computational graph, and use backpropagation over time steps
- Called backpropagation through time (BPTT) algorithm
- Can then apply any general-purpose gradient-based techniques
- Conceptually: first compute the gradients of **the internal nodes**, then compute the gradients of **the parameters**

# Training RNN

- RNN dynamics

$$s^{(t)} = f(x^{(t)}, s^{(t-1)}, \mathbf{W}_s), \quad \mathbf{W}_s = [\mathbf{W} \ \mathbf{U} \ \mathbf{b}]$$

$$o^{(t)} = g(s^{(t)}, \mathbf{V}_o), \quad \mathbf{V}_o = [\mathbf{V} \ \mathbf{c}]$$

- Loss:  $L(\mathbf{x}, \mathbf{y}, \mathbf{W}_s, \mathbf{V}) = \sum_{t=1}^T L^{(t)}(y^{(t)}, o^{(t)})$

- Derivative w.r.t.  $\mathbf{V}_o$

$$\begin{aligned} \frac{\partial L}{\partial \mathbf{V}} &= \frac{1}{T} \sum_{t=1}^T \frac{\partial L^{(t)}(y^{(t)}, o^{(t)})}{\partial \mathbf{V}} \\ &= \frac{1}{T} \sum_{t=1}^T \frac{\partial L^{(t)}(y^{(t)}, o^{(t)})}{\partial o^{(t)}} \frac{\partial g(s^{(t)}, \mathbf{V})}{\partial \mathbf{V}} \end{aligned}$$

**Math formula:**

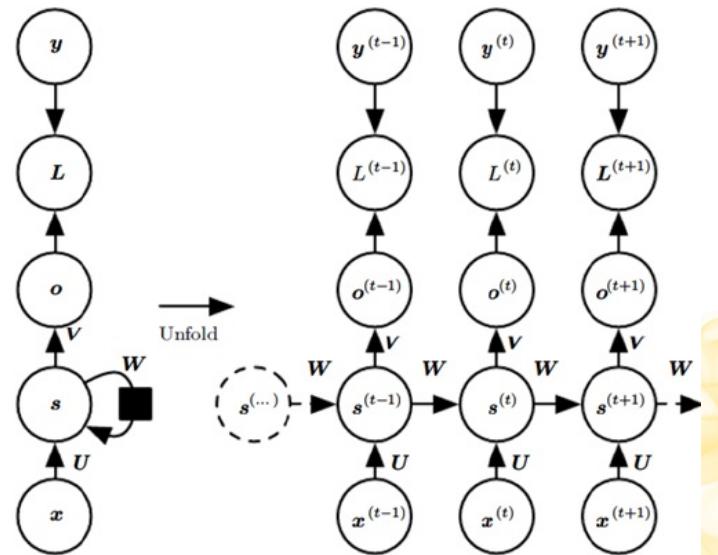
$$a^{(t)} = b + \mathbf{W}s^{(t-1)} + \mathbf{U}x^{(t)}$$

$$s^{(t)} = \tanh(a^{(t)})$$

$$o^{(t)} = c + \mathbf{V}s^{(t)}$$

$$\hat{y}^{(t)} = \text{softmax}(o^{(t)})$$

state transition formula



# Training RNN

- RNN dynamics

$$s^{(t)} = f(x^{(t)}, s^{(t-1)}, \mathbf{W}_s), \quad \mathbf{W}_s = [\mathbf{W} \ \mathbf{U}]$$

$$o^{(t)} = g(s^{(t)}, \mathbf{V})$$

- Loss:**  $L(\mathbf{x}, \mathbf{y}, \mathbf{W}_s, \mathbf{V}) = \sum_{t=1}^T L^{(t)}(y^{(t)}, o^{(t)})$
- Derivatives w.r.t.  $\mathbf{W}_s$

$$\frac{\partial L}{\partial \mathbf{W}_s} = \frac{1}{T} \sum_{t=1}^T \frac{\partial L^{(t)}(y^{(t)}, o^{(t)})}{\partial \mathbf{W}_s}$$

$$= \frac{1}{T} \sum_{t=1}^T \left[ \frac{\partial L^{(t)}(y^{(t)}, o^{(t)})}{\partial o^{(t)}} \frac{\partial g(s^{(t)}, \mathbf{V})}{\partial s^{(t)}} \frac{\partial s^{(t)}}{\partial \mathbf{W}_s} \right]$$

Need to recurrently  
Compute the effect

Can be computed easily

Math formula:

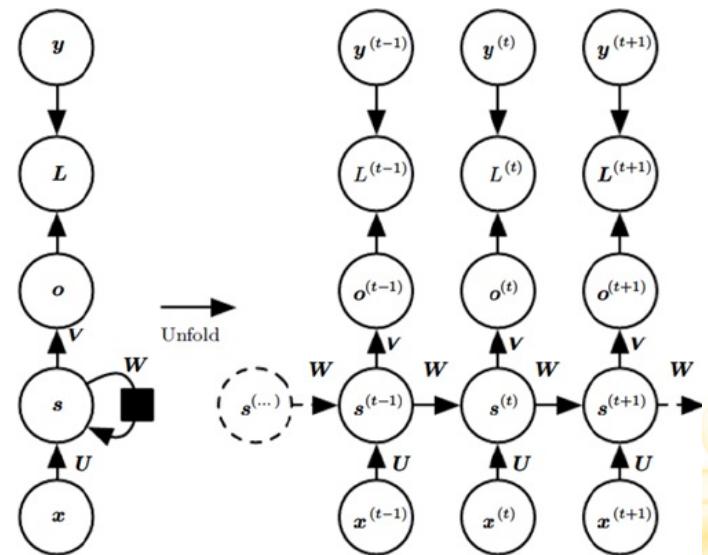
$$a^{(t)} = b + Ws^{(t-1)} + Ux^{(t)}$$

$$s^{(t)} = \tanh(a^{(t)})$$

$$o^{(t)} = c + Vs^{(t)}$$

$$\hat{y}^{(t)} = \text{softmax}(o^{(t)})$$

state transition  
formula



# Training RNN

- RNN dynamics

$$s^{(t)} = f(x^{(t)}, s^{(t-1)}, \mathbf{W}_s), \quad \mathbf{W}_s = [\mathbf{W} \ \mathbf{U}]$$

$$\begin{aligned}\frac{\partial s^{(t)}}{\partial \mathbf{W}_s} &= \frac{\partial f(x^{(t)}, s^{(t-1)}, \mathbf{W}_s)}{\partial \mathbf{W}_s} + \frac{\partial f(x^{(t)}, s^{(t-1)}, \mathbf{W}_s)}{\partial s^{(t-1)}} \frac{\partial s^{(t-1)}}{\partial \mathbf{W}_s} \\ &= \frac{\partial f(x^{(t)}, s^{(t-1)}, \mathbf{W}_s)}{\partial \mathbf{W}_s} + \sum_{\tau=1}^{t-1} \left( \prod_{t'=\tau+1}^t \frac{\partial f(x^{(t')}, s^{(t'-1)}, \mathbf{W}_s)}{\partial s^{(t'-1)}} \right) \frac{\partial f(x^{(\tau)}, s^{(\tau-1)}, \mathbf{W}_s)}{\partial \mathbf{W}_s} \\ &\quad \text{Often truncate this term in practice}\end{aligned}$$

recursion

Math formula:

$$\begin{aligned}a^{(t)} &= b + \mathbf{W}s^{(t-1)} + \mathbf{U}x^{(t)} \\ s^{(t)} &= \tanh(a^{(t)}) \\ o^{(t)} &= c + \mathbf{V}s^{(t)} \\ \hat{y}^{(t)} &= \text{softmax}(o^{(t)})\end{aligned}$$

} state transition formula

# Recurrent Neural Networks

- Use **the same** computational function and parameters across different time steps of the sequence
  - Reduced complexity for improved generalization
- Each time step: takes the input entry and **the previous hidden state** to compute the output entry
- Loss: typically computed at every time step
- Many variants
  - Information about the past can be in many other forms
  - Only output at the end of the sequence

# RNN Variations

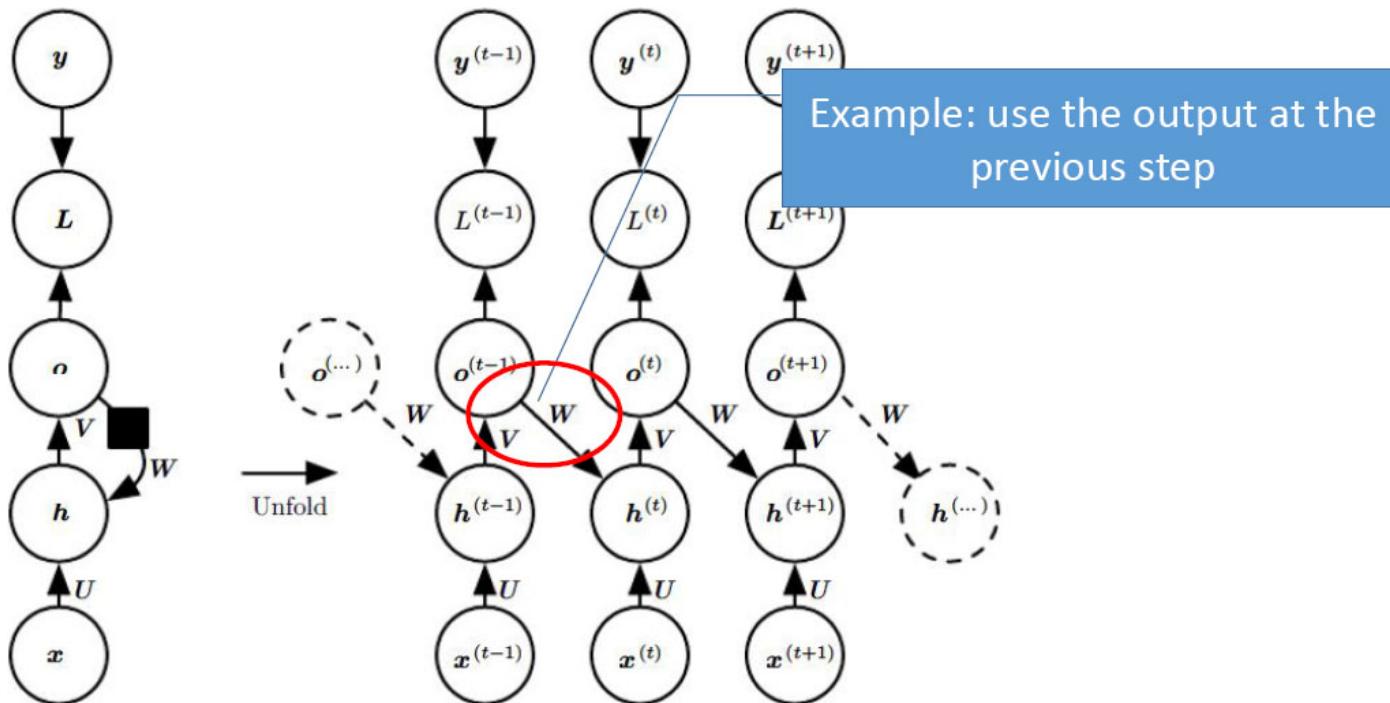


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

# RNN Variations

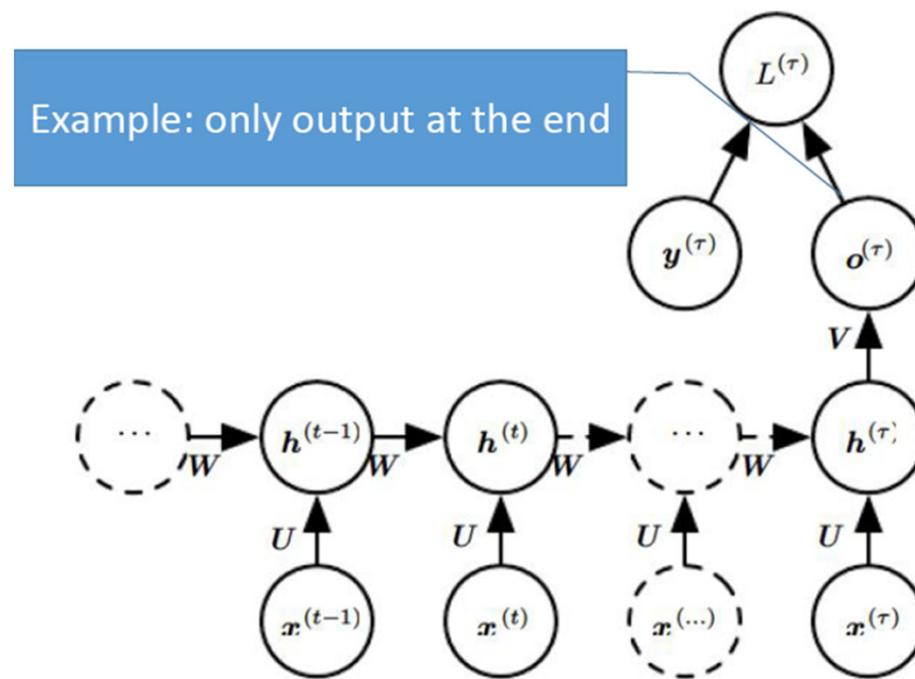


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

# LSTM & GRU



SCIENCE  
ACADEMY

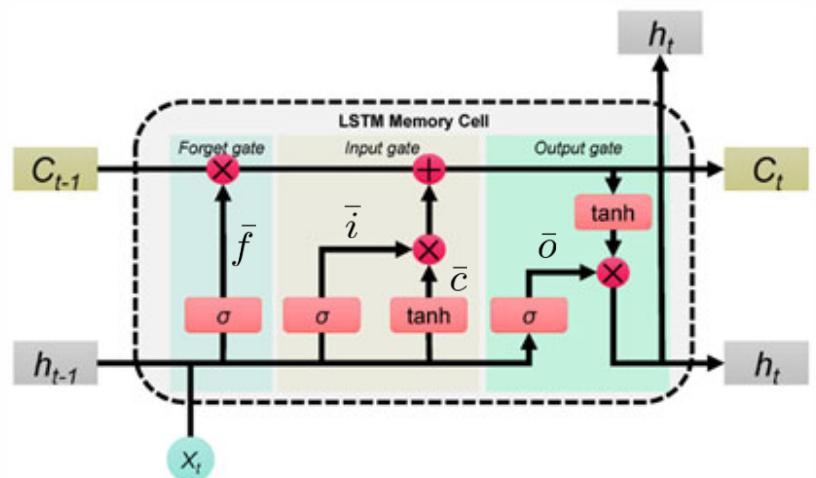
© 2020, UNIVERSITY OF MARYLAND

# Shortcoming of RNNs

- Training a RNN generally a difficult task (cannot be done in parallel due to the sequential nature of data)
- Short memory
  - Called long-term dependency problem
  - Caused by vanishing and exploding gradient problems
- Handled by Long Short-Term Memory (LSTM) or Gated Recurrent Unit (GRU)
  - Introduces memory cells and gates

# Long Short-Term Memory (LSTM)

- Enhancement of RNN architecture
  - Change how hidden states are propagated with the help of additional hidden vector called the “**cell state**”
- Cell state maintains long-term memory
  - Cell states in different temporal layers share more similarity through long-term memory
- Introduces four (4) intermediate vector variables
  - $\bar{i}$  - input variable
  - $\bar{f}$  - forget variable
  - $\bar{o}$  - output variable



# Long Short-Term Memory (LSTM)

- Cell state updated in two steps

- Intermediate variables

- Input gate:  $\bar{i}$
- Forget gate:  $\bar{f}$
- Output gate:  $\bar{o}$
- New C-state:  $\bar{c}$

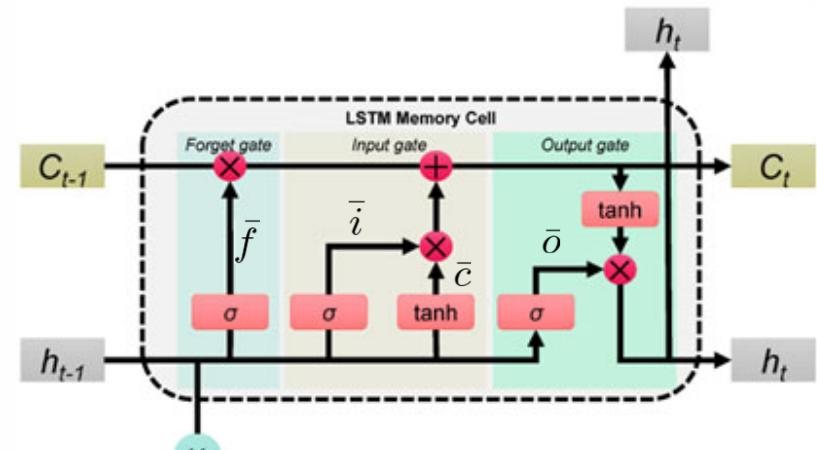
$$\begin{bmatrix} \bar{i} \\ \bar{f} \\ \bar{o} \\ \bar{c} \end{bmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} [\mathbf{W}_h \ \mathbf{W}_x] \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix}$$

- Update of cell state and hidden state

$$c_t = \bar{f} \odot c_{t-1} + \bar{i} \odot \bar{c} \quad [\text{Selectively forget and add to long-term memory (cell state)}]$$

$$h_t = \bar{o} \odot \tanh(c_t) \quad [\text{Selectively leak long-term memory to hidden state}]$$

○ – element-wise product

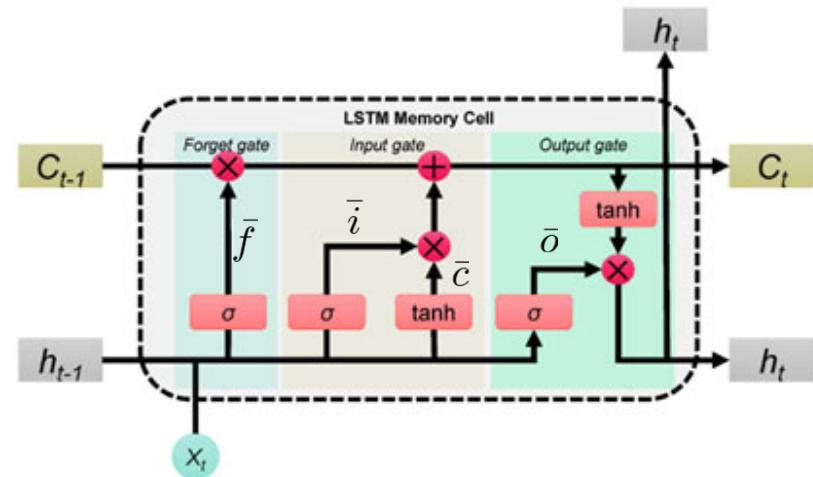


# Long Short-Term Memory (LSTM)

- Update of cell state and hidden state

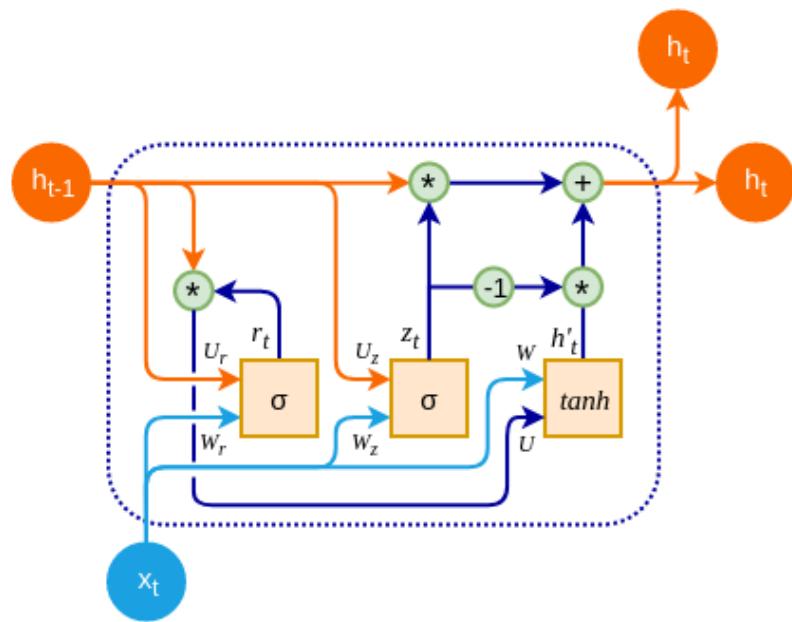
$$c_t = \bar{f} \odot c_{t-1} + \bar{i} \odot \bar{c} \quad [\text{Selectively forget and add to long-term memory (cell state)}]$$
$$h_t = \bar{o} \odot \tanh(c_t) \quad [\text{Selectively leak long-term memory to hidden state}]$$

- Input gate:
  - Forget gate:
  - Output gate:
  - New C-state:
- $$\begin{bmatrix} \bar{i} \\ \bar{f} \\ \bar{o} \\ \bar{c} \end{bmatrix} = \begin{pmatrix} \sigma & & & \\ & \sigma & & \\ & & \sigma & \\ & & & \tanh \end{pmatrix} [\mathbf{W}_h \ \mathbf{W}_x] \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix}$$
- $\bar{c}$  contains newly proposed contents of cell state



# Gated Recurrent Unit (GRU)

- Can be considered a simplification of LSTM without explicit cell state
  - Uses a single “reset” gate to achieve the functions of “forget” and “output” gates in LSTM
- Hidden state updated in two steps
  - Update gate:  $\begin{bmatrix} \bar{z} \\ \bar{r} \end{bmatrix} = \begin{pmatrix} \sigma \\ \sigma \end{pmatrix} [\mathbf{W}_h \mathbf{W}_x] \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix}$
  - Reset gate:  $\begin{bmatrix} r_t \\ z_t \end{bmatrix} = \begin{pmatrix} \sigma \\ \sigma \end{pmatrix} [U_r \quad U_z] \begin{bmatrix} h_{t-1} \\ \bar{r} \odot h_{t-1} \end{bmatrix}$
  - Hidden state update
$$h_t = \bar{z} \odot h_{t-1} + (1 - \bar{z}) \odot \tanh \left( \mathbf{V} \begin{bmatrix} x_t \\ r_t \odot h_{t-1} \end{bmatrix} \right)$$
    - Reset gate determines how much of hidden state to carry over
    - Update gate decides the relative strength of contributions of hidden state



# Bidirectional RNNs

- Many applications: output at time  $t$  may depend on the whole input sequence
- Example in speech recognition: correct interpretation of the current sound may depend on the next few phonemes, potentially even the next few words
- Bidirectional RNNs are introduced to address this

# BiRNNs

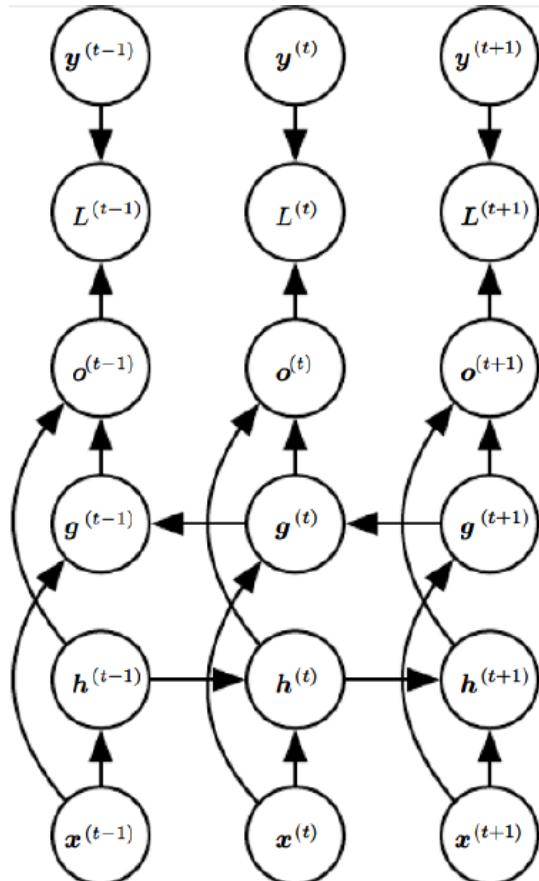


Image source: *Deep Learning*, Goodfellow, Bengio and Courville