

CLUSTERING

DATA/MSML 603: Principles of Machine Learning

Last Time

- Autoencoders
 - Often used for representation learning or dimensionality reduction
- Generative Adversarial Networks (GANs)
 - Consists of two neural networks – one for Generator and the other for Discriminator
 - Generative model used to produce new data points/samples
- Variational Autoencoders
 - Encoder produces a probability distribution (Gaussian)
 - Used to generate new data points/samples
- Transformer Networks
 - Employs self attention, multi-head attention and positional encoding

Clustering





SCIENCE
ACADEMY

© 2020, UNIVERSITY OF MARYLAND

scienceacademy.umd.edu

What Is Clustering?

- Organization of **unlabeled** data into similar groups called **clusters**
 - **Unsupervised** learning
 - A cluster is a collection of data items which are “similar” to each other, and “dissimilar” to data items in other clusters
 - Helps discover patterns, outliers and structure in data and reduces complexity and dimensionality of data
 - Pre-process data before applying a dimensionality reduction technique
 - Clustering-based dimensionality reduction technique





How is Clustering used in Dimensionality Reduction?

- Pre-processing data with clustering
 - Can help remove outliers, noise and irrelevant features
 - Reduce the number of data points by using cluster representatives or centroids (centers of clusters)
 - Can improve the efficiency and effectiveness of dimensionality reduction technique as well as scalability of algorithms
 - Example: Cluster data points into k clusters (for instance, using k-means clustering) and use the k cluster centroids as input for Principal Component Analysis (PCA)
- Clustering-based dimensionality reduction techniques ([topic of next lecture](#))
 - Combine clustering and dimensionality reduction in one step
 - Find a low(er)-dimensional representation of data which preserves cluster structure
 - Example: Spectral clustering uses eigenvectors of a similarity matrix to cluster data

Historic Application of Clustering


- John Snow, a British physician plotted the location of cholera deaths on a map during an outbreak in the 1850s
- The locations indicated that cases were clustered around certain intersections where there were polluted wells—thus exposing both the (source of the) problem and the solution



From: Nina Mishra HP Labs

Computer Vision Application: Image Segmentation

- Pixels with similar characteristics (e.g., color, intensity) grouped into clusters
 - Distinct segments of an image (e.g., objects or regions) represented as clusters
 - Boundary detection



From: Image Segmentation by Nested Cuts, O. Veksler, CVPR2000



Other Applications of Clustering

- Customer segmentation and targeted advertising
 - Study purchase patterns and content recommendation (e.g., streaming)
- Epidemiological analysis
 - Identify clusters of regions or populations with high incidence
- Social media
 - Content personalization
- City planning
 - Resource allocation (playgrounds, schools, parks)
 - Transportation planning based on traffic patterns
- Biology and genetic research

What Do We Need for Clustering?

1. Proximity measure, either:

- Similarity measure $s(x_i, x_k)$: large if x_i and x_k similar
- Dissimilarity (or distance) measure $d(x_i, x_k)$: small if x_i and x_k are similar



2. Criterion function to evaluate clustering


3. Algorithm to compute clustering

- For example, by optimizing the criterion function

1. Distance (Dissimilarity) Measures


- Euclidean distance

- Translation invariant: $d_E(\mathbf{x}^1, \mathbf{x}^2) = \sqrt{\sum_{k=1}^d (x_k^1 - x_k^2)^2} = \|\mathbf{x}^1 - \mathbf{x}^2\|_2$



- Manhattan (city block) distance:

$$d_M(\mathbf{x}^1, \mathbf{x}^2) = \sum_{k=1}^d |x_k^1 - x_k^2| = \|\mathbf{x}^1 - \mathbf{x}^2\|_1$$




- Approximation to Euclidean distance, cheaper to compute
- They are special cases of Minkowski distance

$$d_{MI}(\mathbf{x}^1, \mathbf{x}^2) = \left(\sum_{k=1}^d |x_k^1 - x_k^2|^p \right)^{1/p} = \|\mathbf{x}^1 - \mathbf{x}^2\|_p$$

- p is a positive integer:

2. Cluster Evaluation (A Hard Problem)

- **Intra-cluster cohesion** (compactness):
 - Cohesion measures how close the data points in a cluster are to the cluster centroid
 - Sum of squared error (SSE) is a commonly used measure
- **Inter-cluster separation** (isolation):
 - Different cluster centroids should be far away from one another
 - In most applications, expert judgments are still the key




2. How Many Clusters?

- Possible approaches
 1. Fix the number of clusters to k
 2. Find the best clustering according to the criterion function, e.g., information criterion, such as Akaike information criterion
(number of clusters may vary)




3. Clustering Techniques (1/4)




3. Clustering Techniques (2/4)

- **Hierarchical algorithms** find successive clusters using previously established clusters. These algorithms can be either agglomerative (“bottom-up”) or divisive (“top-down”):
 - **Agglomerative algorithms** begin with each element as a separate cluster and merge them into successively larger clusters
 - **Divisive algorithms** begin with the whole set and proceed to divide it into successively smaller clusters




3. Clustering Techniques (3/4)

- **Partitional algorithms** typically determine all clusters at once, but can also be used as divisive algorithms in the hierarchical clustering
- **Bayesian algorithms** try to generate a posterior distribution over the collection of all partitions of the data



3. Clustering Techniques (4/4)



k-Means Clustering



SCIENCE
ACADEMY

© 2020, UNIVERSITY OF MARYLAND

scienceacademy.umd.edu

k-Means Clustering

- **k-means clustering** (MacQueen, 1967) is a **partitional clustering algorithm**
 - Idea itself goes back further to Steinhaus (1956)
- Let $D = \{\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n\}$ be the set of data points
 - where $\mathbf{x}^i = (x_1^i, x_2^i, \dots, x_r^i)$ is a vector in $X \subset \mathbb{R}^r$, and r is the dimension of data points
- k -means algorithm partitions the given data points into k clusters:
 - k is specified by the user
 - Each cluster has a **cluster center**, called the **centroid**

k-Means Algorithm

- Given k , the ***k-means algorithm*** works as follows:
 - Choose k (random) data points (seeds) to be the initial centroids, i.e., cluster centers
 - Assign each data point to the closest centroid
 - Re-compute the centroids using the current cluster memberships
 - If a convergence criterion is not met, repeat steps 2 and 3

k-Means Convergence (Stopping) Criterion

- No (or minimum) re-assignments of data points to different clusters, or
- No (or minimum) change of centroids, or
- minimum decrease in the **sum of squared error (SSE)**,


$$SSE = \sum_{j=1}^k \sum_{\mathbf{x} \in C_j} d(\mathbf{x}, \mathbf{m}_j)^2$$

- C_j is the j th cluster,
- \mathbf{m}_j is the centroid of cluster C_j (the mean vector of all the data points in C_j),
- $d(\mathbf{x}, \mathbf{m}_j)$ is the (Euclidian) distance between data point \mathbf{x} and centroid \mathbf{m}_j .

$$\mathbf{m}_j = \frac{1}{|C_j|} \sum_{\mathbf{x} \in C_j} \mathbf{x}$$

k-Means Clustering Example (1/6)

- Step 1 Randomly initialize the cluster centers (synaptic weights)




$(k = 3)$

k-Means Clustering Example (2/6)

- Step 2


Determine cluster membership for each input
("winner-takes-all" inhibitory circuit)



k-Means Clustering Example (3/6)

- Step 3


Re-estimate cluster centers (adapt synaptic weights)



$$\mathbf{c}_j = \frac{1}{|C_j|} \sum_{\mathbf{x} \in C_j} \mathbf{x}$$


k-Means Clustering Example (4/6)

Result of first iteration




k-Means Clustering Example (5/6)

Second iteration



$$\mathbf{c}_j = \frac{1}{|C_j|} \sum_{\mathbf{x} \in C_j} \mathbf{x}$$

k-Means Clustering Example (6/6)



Why Use k-Means?


- Advantages:
 - **Simple:** easy to understand and to implement
 - **Efficient:** Time complexity: $O(tkn)$, where n is the number of data points, k is the number of clusters, and t is the number of iterations
 - Since both k and t are small, k -means is considered a linear algorithm
- k -means is the most popular clustering algorithm
- Comment:
 - It terminates at a local optimum if SSE is used
 - The global optimum is hard to find due to complexity

Weaknesses of k-Means

- The algorithm is only applicable if the mean is defined
 - For categorical data, k -modes clustering is used - the centroid is represented by the most frequent values (called the mode of cluster)
- The user needs to specify k
 - Need to know approximately how many clusters to look for
- The algorithm is sensitive to **outliers**
 - Outliers are data points that are very far away from other data points
 - Outliers could be errors in the data recording or some special data points with very different values

Outliers

- Undesirable clusters



$k = 2$

- Ideal clusters




Dealing With Outliers


- Remove some data points that are much further away from the centroids than other data points
 - To be safe, we may want to monitor these possible outliers over a few iterations and then decide to remove them
- Perform random sampling: by choosing a small subset of the data points, the chance of selecting an outlier is much smaller
 - Assign the rest of the data points to the clusters by distance or similarity comparison, or classification

Sensitivity To Initial Seeds


- Sensitive to initial seeds




Random selection of seeds (centroids)



Iteration 1




Iteration 2



Random selection of seeds (centroids)




Iteration 1




Iteration 2

Special Data Structures

- The k -means algorithm is not suitable for discovering clusters that are not hyper-ellipsoids (or hyper-spheres)



(A): Two natural clusters




(B): k -means clusters

k-Means Summary


- Despite weaknesses, *k*-means clustering is still the most popular algorithm due to its simplicity and efficiency
- No clear evidence that any other clustering algorithm performs better in general
- Comparing different clustering algorithms is a difficult task
 - No one knows the correct clusters!

Application to Visual Object Recognition: Dictionary Learning (Bag of Words)

- Dictionary learning – find a set of basic image features (in **dictionary**)
 - Group similar image patches together into a cluster to capture unique characteristics of different image regions
 - Dictionary learns to represent features specific to each cluster using a set of dictionary **atoms**
 - Dictionary is a matrix containing atoms (basic features) as columns
 - An image written as a linear combination of atoms (usually interested in sparse representation)




Learning the Visual Vocabulary (1/2)



Learning the Visual Vocabulary (2/2)

- Extract features from different patches of images
- Extracted features clustered into a set of “visual words” (common visual patterns)



k-Medoids Clustering


k-Medoids

- Almost the same as k -means
- Difference
 - k -means: centers are calculated (centroids)
 - k -medoids: centers (medoids) are selected from samples
 - Data point that minimizes the sum distances to all other points in the cluster
 - Distance metric: k -means (L2 norm), k -medoids (L1 norm)

$$\mathbf{m}_j = \frac{1}{|C_j|} \sum_{\mathbf{x} \in C_j} \mathbf{x}$$


Outliers

- Advantage
 - k -medoids is more robust to noise and outliers than k -means




Hierarchical Clustering

Clustering Techniques




Hierarchical Clustering


- Up to now, we considered “flat” clustering:



- For some data, hierarchical clustering is more appropriate than “flat” clustering
- Hierarchical clustering



Example: Biological Taxonomy



Dendrogram

- Preferred way to represent hierarchical clustering (a tree diagram)




Diagram that
shows the
arrangements
of the clusters

Types of Hierarchical Clustering


- **Divisive (top down) clustering**

- Start with all data points in one cluster, the root, then
 - Split the root into a set of child clusters
 - Each child cluster is recursively divided further
 - Stop when only singleton clusters of individual data points remain, i.e., each cluster with only a single point

- **Agglomerative (bottom up) clustering**

- The dendrogram is built from the bottom level by
 - merging the most similar (or nearest) pair of clusters
 - stopping when all the data points are merged into a single cluster (i.e., the root cluster).

Divisive Hierarchical Clustering




- Any “flat” algorithm which produces a fixed number of clusters can be used
 - Set $c=2$

Agglomerative Hierarchical Clustering

initialize with each example in singleton cluster

while there is more than 1 cluster

1. find 2 nearest clusters
2. merge them



Four common ways to measure cluster distance

1. minimum distance $d_{\min}(D_i, D_j) = \min_{x \in D_i, y \in D_j} \|x - y\|$

2. maximum distance $d_{\max}(D_i, D_j) = \max_{x \in D_i, y \in D_j} \|x - y\|$

3. average distance $d_{avg}(D_i, D_j) = \frac{1}{n_i n_j} \sum_{x \in D_i} \sum_{y \in D_j} \|x - y\|$


4. mean distance $d_{mean}(D_i, D_j) = \|\mu_i - \mu_j\|$

Divisive vs. Agglomerative

- **Agglomerative** clustering is faster to compute, in general
- **Divisive** clustering may be less “blind” to the global structure of the data

Divisive

when taking the first step (split), have access to all the data; can find the best possible split in 2 parts



Agglomerative


when taking the first step merging, do not consider the global structure of the data, only look at pairwise structure



Gaussian Mixture Model


The Evils of “Hard Assignments”

- Clusters may overlap
- Some clusters may be “wider” than others
- Distances can be deceiving!
- Aggregate model much richer than component models (for individual clusters)



Probabilistic Clustering


- Try a **probabilistic model!**
 - Soft clustering
 - Allows overlaps, clusters of different sizes, etc.
- Can tell a generative story for data
 - $\mathbf{P}(X, Y) = \mathbf{P}(X \mid Y)\mathbf{P}(Y)$
- **Challenge:** We need to estimate model parameters **without labels**



Y	X ₁	X ₂
??	0.1	2.1
??	0.5	-1.1
??	0.0	3.0
??	-0.1	-2.0
??	0.2	1.5
...

General Gaussian Mixture Model Assumption

- $P(Y)$: There are K components ($K > 1$)
- $P(X|Y)$: Each component k generates data from a multivariate Gaussian with mean μ_k and covariance matrix Σ_k
- Each data point is sampled from a two-step generative process (assuming **known parameters**):
 1. Choose **component** k with probability $P(Y = k)$
 2. Generate datapoint from distribution $N(\mu_k, \Sigma_k)$



Gaussian mixture model (GMM)

What Model Should We Use?


- In general, choice of model depends on X !
- Here, maybe Gaussian Naïve Bayes?
 - Numbers of samples from different components given by a multinomial random variable
 - Given fixed components, X_i (independent) Gaussian

Component probability: $\mathbf{P}(Y_i = k) = \pi_k, k = 1, \dots, K$

Conditional PDF (scalar case):

$$p(x | Y = k) = \frac{1}{\sqrt{2\pi}\sigma_k} \exp\left(-\frac{(x - \mu_k)^2}{2\sigma_k^2}\right)$$

conditional probability density



Y	X ₁	X ₂
??	0.1	2.1
??	0.5	-1.1
??	0.0	3.0
??	-0.1	-2.0
??	0.2	1.5
...

Could We Make Fewer Assumptions?

- Vector case – multivariate Gaussian distribution


Conditional PDF (vector case):

$$p(\mathbf{x}_j \mid Y = k) = \frac{1}{(2\pi)^{m/2} |\Sigma_k|^{0.5}} \exp \left(-\frac{1}{2} (\mathbf{x}_j - \boldsymbol{\mu}_k)^T \Sigma_k^{-1} (\mathbf{x}_j - \boldsymbol{\mu}_k) \right),$$

$\mathbf{x}_j \in \mathbb{R}^m$

↑
conditional probability density



m = dimension of \mathbf{x}_j



Multivariate Gaussians (1/4)

$$X \sim N(\boldsymbol{\mu}, \Sigma) : p(\mathbf{x}_j) = \frac{1}{(2\pi)^{m/2} |\Sigma|^{0.5}} \exp \left(-\frac{1}{2} (\mathbf{x}_j - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x}_j - \boldsymbol{\mu}) \right), \quad |\Sigma| = \det(\Sigma)$$

m = dimension of \mathbf{x}



$\Sigma \propto$ identity matrix





SCIENCE
ACADEMY

© 2020, UNIVERSITY OF MARYLAND

DATA/MSML 603

Multivariate Gaussians (2/4)



$$p(\mathbf{x}_j) = \frac{1}{(2\pi)^{m/2}|\Sigma|^{0.5}} \exp\left(-\frac{1}{2}(\mathbf{x}_j - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x}_j - \boldsymbol{\mu})\right)$$



$\Sigma = \text{diagonal matrix}$
 X_i are independent *ala* Gaussian NB


Multivariate Gaussians (3/4)

$$p(\mathbf{x}_j) = \frac{1}{(2\pi)^{m/2} |\Sigma|^{0.5}} \exp\left(-\frac{1}{2} (\mathbf{x}_j - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x}_j - \boldsymbol{\mu})\right) =: N(\mathbf{x}_j; \boldsymbol{\mu}, \Sigma)$$



Σ = arbitrary (semidefinite) matrix:
- specifies rotation (change of basis)
- eigenvalues specify relative elongation


Multivariate Gaussians (4/4)



$$p(\mathbf{x}_j) = \frac{1}{(2\pi)^{m/2} |\Sigma|^{0.5}} \exp \left(-\frac{1}{2} (\mathbf{x}_j - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x}_j - \boldsymbol{\mu}) \right)$$


Mixtures of Gaussians (1/4)

Old Faithful Data Set




Mixtures of Gaussians (2/4)

Old Faithful Data Set



Single Gaussian



Mixture of two Gaussians


Mixtures of Gaussians (3/4)

- Allows us to combine simple models into a complex model




$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

↑
Component
Mixing coefficient

$$\forall k : \pi_k \geq 0 \quad \sum_{k=1}^K \pi_k = 1$$




Mixtures of Gaussians (4/4)



Eliminating Hard Assignments to Clusters (1/2)


- Model data as mixture of multivariate Gaussians



(with known
components)


Eliminating Hard Assignments to Clusters (2/2)

- Model data as mixture of multivariate Gaussians



Eliminating Hard Assignments to Clusters

- Model data as mixture of multivariate Gaussians



Shown is the *posterior probability* that a point was generated from i^{th} Gaussian: $\mathbf{P}(Y = i \mid \mathbf{x})$

ML Estimation in Supervised Settings

- Recall: **Univariate Gaussian**

$$\mu_{MLE} = \frac{1}{N} \sum_{i=1}^N x_i \quad \sigma_{MLE}^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \hat{\mu})^2$$

- **Mixture of multivariate Gaussians:** for each component distribution

ML estimate for each of the Multivariate Gaussians is given by:

$$\mu_{ML}^k = \frac{1}{n} \sum_{j=1}^n x_n \quad \Sigma_{ML}^k = \frac{1}{n} \sum_{j=1}^n (\mathbf{x}_j - \mu_{ML}^k)(\mathbf{x}_j - \mu_{ML}^k)^T$$


Just sums over x generated from the k 'th Gaussian




That was easy! But what if there are unobserved data?

- MLE:
$$=: p(\mathbf{x}_j, y_j)$$
 - $\arg \max_{\boldsymbol{\theta}} \prod_{j=1}^n p(\mathbf{x}_j \mid Y_j = y_j) \mathbf{P}(Y_j = y_j)$
 - $\boldsymbol{\theta} = \{(\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k), k = 1, \dots, K\}$: all model parameters including, class probabilities, means and variances
- Problem: We do not know y_j , i.e., which component is responsible for generating each data point
- Maximize ***marginal likelihood***

$$\arg \max_{\boldsymbol{\theta}} \prod_{j=1}^n p(\mathbf{x}_j) = \arg \max_{\boldsymbol{\theta}} \prod_{j=1}^n \left(\sum_{k=1}^K p(\mathbf{x}_j \mid Y_j = k) \mathbf{P}(Y_j = k) \right)$$



How do we optimize?



- Maximize ***marginal likelihood***

$$\arg \max_{\boldsymbol{\theta}} \prod_{j=1}^n p(\mathbf{x}_j) = \arg \max_{\boldsymbol{\theta}} \prod_{j=1}^n \left(\sum_{k=1}^K p(\mathbf{x}_j \mid Y_j = k) \mathbf{P}(Y_j = k) \right)$$

$= p(\mathbf{x}_j, k)$

- Almost always a hard problem
 - Usually no closed form solution
 - Even when $\log(p(X, Y))$ is convex, $\log(p(X))$ generally is not
 - For all but the simplest $p(X)$, we will have to use gradient ascent, in a big messy space with many local optima

Learning General Mixtures of Gaussian

- Marginal likelihood with multivariate Gaussian:

$$\prod_{j=1}^n p(\mathbf{x}_j) = \prod_{j=1}^n \left(\sum_{k=1}^K \underbrace{\frac{1}{(2\pi)^{m/2} |\Sigma_k|^{0.5}} \exp \left(-\frac{1}{2} (\mathbf{x}_j - \boldsymbol{\mu}_k)^T \Sigma_k^{-1} (\mathbf{x}_j - \boldsymbol{\mu}_k) \right)}_{= p(\mathbf{x}_j \mid Y_j = k)} \cdot \mathbf{P}(Y_j = k) \right)$$

- Need to differentiate and solve for the parameters $\theta = \{(\pi_k, \boldsymbol{\mu}_k, \Sigma_k), k = 1, \dots, K\}$
- There is no closed-form solution, gradient is complex, and many local optima
- How do we solve it then? Is there a better way?***

Expectation Maximization

- Recall (for one data point)

$$\begin{aligned} p(\mathbf{x}_i) &= \sum_{k=1}^K \underbrace{\frac{1}{(2\pi)^{m/2} |\Sigma_k|^{0.5}} \exp\left(-\frac{1}{2} (\mathbf{x}_i - \boldsymbol{\mu}_k)^T \Sigma_k^{-1} (\mathbf{x}_i - \boldsymbol{\mu}_k)\right)}_{= N(\mathbf{x}_i; \boldsymbol{\mu}_k, \Sigma_k)} \cdot \underbrace{\mathbf{P}(Y_i = k)}_{= \pi_k} \\ &= \sum_{k=1}^K \pi_k \cdot N(\mathbf{x}_i; \boldsymbol{\mu}_k, \Sigma_k) \end{aligned}$$

- Log-likelihood: $\ell(\mathbf{x}_1, \dots, \mathbf{x}_n; \boldsymbol{\theta}) = \sum_{i=1}^n \log \left(\sum_{k=1}^K \pi_k \cdot N(\mathbf{x}_i; \boldsymbol{\mu}_k, \Sigma_k) \right)$
- Differentiate w.r.t. each parameter and set the derivative to 0 (assume scalar instead of vector for simplicity)

$$E.g., \text{Derivative w.r.t. } \mu_l : \sum_{i=1}^n \frac{1}{\sum_{k=1}^K \pi_k \cdot N(x_i; \boldsymbol{\mu}_k, \Sigma_k)} \pi_l N(x_i; \mu_l, \sigma_l^2) \frac{(x_i - \mu_l)}{\sigma_l^2} = 0$$

- Cannot analytically solve for μ_l

Expectation Maximization

- Key idea: Break the optimization w.r.t. parameters into two steps and alternate between them
- 1. **Expectation (E) step:** Evaluate the **posterior probabilities** using the current estimates of component Gaussian distributions and mixing components $\{(\hat{\pi}_k, \hat{\mu}_k, \hat{\Sigma}_k), k = 1, \dots, K\}$

$$\mathbf{P}(Y_i = l \mid X_i = \mathbf{x}_i) = \frac{p(\mathbf{x}_i \mid Y_i = l)\mathbf{P}(Y_i = l)}{p(\mathbf{x}_i)} = \frac{\hat{\pi}_l \cdot N(\hat{\mu}_l, \hat{\Sigma}_l)}{\sum_{k=1}^K \hat{\pi}_k \cdot N(\hat{\mu}_k, \hat{\Sigma}_k)} = \gamma_i(l), \quad l = 1, \dots, K$$

- 2. **Maximization (M) step:** Update new parameters $\hat{\theta} = \{(\hat{\pi}_k, \hat{\mu}_k, \hat{\Sigma}_k), k = 1, \dots, K\}$ with fixed posterior probabilities

$$N_l = \sum_{i=1}^n \gamma_i(l), \quad \hat{\mu}_l = \frac{1}{N_l} \sum_{i=1}^n \gamma_i(l) \mathbf{x}_i,$$

$$\hat{\Sigma}_l = \frac{1}{N_l} \sum_{i=1}^n \gamma_i(l) (\mathbf{x}_i - \hat{\mu}_l)(\mathbf{x}_i - \hat{\mu}_l)^T \quad \text{and} \quad \hat{\pi}_l = \frac{N_l}{n}$$

N_l = effective # of data points assigned to component l


Expectation Maximization

Expectation-Maximization (EM) algorithm

1. Initialize parameters $\hat{\theta} = \{(\hat{\pi}_k, \hat{\mu}_k, \hat{\Sigma}_k), k = 1, \dots, K\}$ and evaluate the log-likelihood
 2. **Expectation (E) step:** Evaluate the **posterior probabilities** $\gamma_i(l) = \mathbf{P}(Y_i = l \mid X_i = \mathbf{x}_i)$ using the current estimates of parameters
 3. **Maximization (M) step:** Update new parameters $\hat{\theta} = \{(\hat{\pi}_k, \hat{\mu}_k, \hat{\Sigma}_k), k = 1, \dots, K\}$ with fixed posterior probabilities
 4. Evaluate the log-likelihood with the new parameter estimates. Stop if stopping criterion is met. Otherwise, go back to Step 2
- Useful fact: log-likelihood increases after each iteration


Gaussian Mixture Example (1/8)

- Start with initial parameters




Gaussian Mixture Example (2/8)

- After first iteration




Gaussian Mixture Example (3/8)

- After 2nd iteration




Gaussian Mixture Example (4/8)

- After 3rd iteration




Gaussian Mixture Example (5/8)

- After 4th iteration




Gaussian Mixture Example (6/8)

- After 5th iteration



Gaussian Mixture Example (7/8)

- After 6th iteration



Gaussian Mixture Example (8/8)

- After 20th iteration

